

# My Bachelorthesis title

Bachelorarbeit von Philipp Hinz  
Tag der Einreichung: 19. Februar 2023

1. Gutachten: Gutachter 1
  2. Gutachten: Gutachter 2
  3. Gutachten: noch einer
  4. Gutachten: falls das immernoch nicht reicht
- Darmstadt



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

Fachbereich Informatik  
Institut  
Arbeitsgruppe

---

## **Erklärung zur Abschlussarbeit gemäß § 22 Abs. 7 und § 23 Abs. 7 APB der TU Darmstadt**

---

Hiermit versichere ich, Philipp Hinz, die vorliegende Bachelorarbeit ohne Hilfe Dritter und nur mit den angegebenen Quellen und Hilfsmitteln angefertigt zu haben. Alle Stellen, die Quellen entnommen wurden, sind als solche kenntlich gemacht worden. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen.

Mir ist bekannt, dass im Fall eines Plagiats (§ 38 Abs. 2 APB) ein Täuschungsversuch vorliegt, der dazu führt, dass die Arbeit mit 5,0 bewertet und damit ein Prüfungsversuch verbraucht wird. Abschlussarbeiten dürfen nur einmal wiederholt werden.

Bei der abgegebenen Thesis stimmen die schriftliche und die zur Archivierung eingereichte elektronische Fassung gemäß § 23 Abs. 7 APB überein.

Bei einer Thesis des Fachbereichs Architektur entspricht die eingereichte elektronische Fassung dem vorgestellten Modell und den vorgelegten Plänen.

Darmstadt, 19. Februar 2023

---

P. Hinz



# Inhaltsverzeichnis

---

|          |  |          |
|----------|--|----------|
| <b>1</b> | <b>Introduction</b>                        | <b>4</b> |
| <b>2</b> | <b>Case Study</b>                          | <b>5</b> |
| <b>3</b> | <b>Extendable CmRDT</b>                    | <b>6</b> |
| 3.1      | Overview . . . . .                         | 6        |
| 3.2      | Concepts . . . . .                         | 8        |
| 3.3      | Authentication and authorization . . . . . | 10       |



---

# 1 Introduction

---

---

## 2 Case Study

---

Our case study is named Ratable. With Ratable users can create ratable objects and let a group of people rate them. Ratable is an application, users can interact with on their mobile devices or desktop computers.

The core idea contrary to usual rating services is that not everyone can rate a Ratable. With a newly created Ratable two links are provided. A rate and a view link. Only people with the rate link can rate and view the ratable.

Ratable is implemented as a cloud-native local-first application. Ratable allows users to use the application without an internet connection and provides updates/changes in real-time. This provides for a very fluent usage of Ratable because we do not need to have any loading screens. The only time we need to load something is when we access some Ratable for the first time. Changes to Ratables are done instantly and synchronized when possible.

Especially important for local-first applications is the feature that all trust is given to individual users instead of a server. This means that the server for Ratable plays only a secondary role and does not do much more than distributing state. All authorization and authentication are done by the clients themselves. This allows us to enable end-to-end encryption of the state.

---

## 3 Extendable CmRDT

---

In this chapter, we will look into what ECmRDT's are, what they are trying to do and how they are used to achieve the goal.

---

### 3.1 Overview

---

Extendable CmRDT's are normal CmRDT's with the ability to build extensions for them. Here CmRDT's are operation based CRDT's. Extensions are able to verify and mutate events and state. They can also interact with other extensions.

With ECmRDT's we want to enable developers to create extensions that can be enabled for data types to easily extend them with additional functionality. This functionality can be for example something like verification of replicaId's in events, change of some variables in state before or after applying an event, migrations of events or authentication and authorization of events. Especially the last example is the primary reason for the usage of ECmRDT's. Contrary to normal CRDT's they allow us to build in security.

Generally ECmRDT's work with events and state. In our design we do only use event sourcing at server side. That means that an ECmRDT itself only contains state. Users can then create an event and apply it to the ECmRDT and distribute it to other clients. Because events are signed, they can not be forged. How an event is applied to state in detail depends on and is defined in the domain logic. Extensions now come into place directly before applying the event to the state. The extensions are defined in something like a pipeline and can either forward the event to the next extension or fail. Thereby events can either change the data coming in or coming out.

Lets say we want to implement a counter that can only be incremented by the owner. We would first define the state as a integer counter and one event that increments the state counter by one. The context and generally how these concepts relate to each other will be covered in the next section.

```
case class Counter(  
  val value: Int,  
)  
  
case class CounterContext(  
  val replicaId: ReplicaId,  
) extends IdentityContext  
  
sealed trait CounterEvent extends Event[Counter, CounterContext]  
  
case class AddCounterEvent() extends CounterEvent:
```

---

```
def asEffect =
  (state, context, meta) => EitherT.pure(
    state.copy(value = state.value + 1)
  )
```

Then we would enable an build-in extensions that only lets events pass when they are send by the owner of the ratable.

```
object Counter:
  given EffectPipeline[Counter, CounterContext] = EffectPipeline(
    SingleOwnerEffectPipeline()
  )
```

Using only a few lines of code we have now a secure incrementable counter that can only be incremented by the owner. This counter can now be used like the following. The example will first create a replicaId and an initial state. With both we can create a event and apply it to our created state. In the end we print the state before and after applying our event.

```
def main(using Crypt) =
  for
    // Step 1: Create replicaId
    replicaId <- EitherT.liftF(PrivateReplicaId())

    // Step 2: Create initial state.
    counter = ECmRDT[Counter, CounterContext, CounterEvent](Counter(0))

    // Step 3: Create event.
    eventPrepared = counter.prepare(
      AddCounterEvent(),
      CounterContext(replicaId)
    )

    // Step 4: Verify and advance state.
    newCounter <- counter.effect(eventPrepared, MetaContext(
      AggregateId.singleton(replicaId),
      replicaId
    ))

  yield
    println(s"Old counter: ${counter.state.value}") // 0
    println(s"New counter: ${newCounter.state.value}") // 1
```

How exactly the ECmRDT's and the SingleOwnerEffectPipeline work will be in the next section.

---

## 3.2 Concepts

---

In this section, we will introduce concepts of our ECmRDT with definitions, code examples and also explain how certain relate to each other.

### ReplicaId

The ReplicaId is a unique identifier for a user. The ReplicaId consists of a public/private key pair. With this we can sign and verify events send by a user. Because the private key only exists in the corresponding replica as a privateReplicaId, the ReplicaId as a data structure only contains the public key for identification. More about the usage of the ReplicaId is covered in the chapter about authentication and authorization.

```
case class ReplicaId(  
  val publicKey: BinaryData  
)
```

### AggregateId

The AggregateId is a unique identifier for an ECmRDT/aggregate. The AggregateId is a combination of a ReplicaId and random bytes. The ReplicaId inside the aggregateId represents the owner of the aggregate. We need to store the ReplicaId to prevent other replica's of creating aggregate with the same AggregateId maliciously. The random bytes are used to prevent collisions of AggregateIds within a group of aggregates of a replica.

```
case class AggregateId(  
  val replicaId: ReplicaId,  
  val randomBytes: BinaryData  
)
```

### Effect

An Effect is a function that takes a state, an context, an MetaContext and returns a new state or an RatableError in future. By returning an RatableError we can verify the Event together with the given parameters. The effect is an asynchronous operation because we need to use cryptograpic operations to verify the Event which are usally implemented asynchronously.

```
type Effect[A, C] = (A, C, MetaContext) => EitherT[Future, RatableError, A]
```

### Event

If the user wants to take any actions on the ECmRDT, he has to create an event. Every event is associated with a context. An Event is applied to the ECmRDT by converting it into an Effect.

```
trait Event[A, C]:  
  def asEffect: Effect[A, C]
```



---

## MetaContext

The MetaContext is a context that is passed to the ECmRDT when applying an Event. It contains information about the ownerReplicaId and the aggregateId.

The need results from the fact, that the initialization of ECmRDTs also happens through an Event. While processing the first Event we do not have any state yet. Therefore we could normally not validate the Event. The MetaContext is mainly used to validate the first Event by checking if the Event was created by the ownerReplicaId.

```
case class MetaContext(  
  val aggregateId: AggregateId,  
  val ownerReplicaId: ReplicaId,  
)
```

## ECmRDTEventWrapper

The ECmRDTEventWrapper is a wrapper for an Event and a context. It is used to bundle the Event with the context and the clock of the ECmRDT when the Event was created. The clock is used to prevent duplications of Events.

```
case class ECmRDTEventWrapper[A, C, +E <: Event[A, C]](  
  val time: Long,  
  val event: E,  
  val context: C,  
)
```

## Context

The context contains meta information about a single Event but is the same for all Events of a single ECmRDT implementation. The idea is to compose the context out of multiple traits.

An core context trait is the IdentityContext. This context trait is always required and contains the ReplicaId of the creator of this Event. The validation of this identity is not done by the ECmRDT but must be done externally. More on this in the section about authentication and authorization.

```
trait IdentityContext:  
  def replicaId: ReplicaId
```

## ECmRDT

Our core concept is the ECmRDT. The ECmRDT consists of a state and a clock. The state is the actual data of the ECmRDT. The clock is a VectorClock to prevent duplications of Events. Our ECmRDT does not handle distribution of Events nor storing pending Events. This has to be done by the user of the ECmRDT.

Our ECmRDT supports two operations. One to prepare an Event and one to apply an prepared event called ECmRDTEventWrapper. The prepare operation is used to aquire additional information from the ECmRDT, specifically the clock and bundle the event with an context into one ECmRDTEventWrapper. The apply operation is used to apply the Event to the ECmRDT.

---

```
case class ECmRDT[A, C <: IdentityContext, E <: Event[A, C]](
  val state: A,
  val clock: VectorClock = VectorClock(Map.empty)
):
  def prepare(
    event: E, context: C
  )(
    using effectPipeline: EffectPipeline[A, C]
  ): ECmRDTEventWrapper[A, C, E] = ...

  def effect(
    wrapper: ECmRDTEventWrapper[A, C, E], meta: MetaContext
  )(
    using effectPipeline: EffectPipeline[A, C]
  ): EitherT[Future, RatableError, ECmRDT[A, C, E]] = ...
```

### EffectPipeline

The EffectPipeline is a simple structure that converts an existing Effect into a new Effect. It is the core piece in implementing extensions. The EffectPipeline is implicitly specified by a ECmRDT implementation and is used to implement additional functionality like logging, validation or mutation. Functionality provided by the EffectPipeline is applied to all Events of a ECmRDT.

```
trait EffectPipeline[A, C]:
  def apply(effect: Effect[A, C]): Effect[A, C]
```

---

## 3.3 Authentication and authorization

---